

eBPF网络包处理：实验手册

基于eBPF的防火墙

一、实验环境

- VMware Workstation Pro
- Ubuntu 22.04

本次实验的相关依赖应当与Ubuntu版本无关，但仅保证Ubuntu 22.04可以正常执行，如果遇到问题可以联系助教。

二、准备工作

1. 安装相关依赖

在终端执行以下命令：

```
sudo apt install clang llvm libelf-dev libbpf-dev libpcap-dev gcc-multilib build-essential linux-tools-$(uname -r)
```

三、实验介绍

1. 实验简述

上上周我们利用内核模块实现了一个简单的防火墙，本次实验中，我们会利用eBPF来实现相同的目标，包括：

- 基于eBPF/TC，屏蔽由223.5.5.5发送的数据包
- 基于eBPF/XDP，屏蔽由223.5.5.5发送的数据包

2. 实验思路

与内核模块的实验类似，eBPF下的思路如下：

1. 分析与网络数据相关的xdp_md或__sk_buff结构体
2. 判断相应结构体数据是否满足需求，并返回对应的响应参数

四、开始实验

1. 解析器函数parser

在内核模块中，我们使用了内核提供的ip_hdr()函数来对sk_buff结构体进行解析，同时取出IP报文的头部。

但在eBPF中，由于eBPF虚拟机的存在以及内核安全性的要求等，我们没有办法直接调用ip_hdr()函数来提取IP头部，如果强行使用，可以正常编译，但在尝试加载时会报错：

```
zhaidoudou123@Lab-virtual-machine:~/eBPF/packet/tc$ sudo tc filter add dev ens33 egress bpf direct-action obj tc.o sec tc
libbpf: failed to find BTF for extern 'ip_hdr': -3
ERROR: opening BPF object file failed
Unable to load program
```

幸运的是，XDP项目在其教程xdp-tutorial中提供了一个现成的parser库parsing_helper.c，我们可以直接调用。

我们以IP头部的解析器parse_iphdr为例来简要介绍：

```
/* Header cursor to keep track of current parsing position */
struct hdr_cursor {
    void *pos;
};

static __always_inline int parse_iphdr(struct hdr_cursor *nh,
                                       void *data_end,
                                       struct iphdr **iphdr)
{
    struct iphdr *iph = nh->pos;
    int hdrsize;

    if (iph + 1 > data_end)
        return -1;

    hdrsize = iph->ihl * 4;
    /* Sanity check packet field is valid */
    if (hdrsize < sizeof(*iph))
        return -1;

    /* Variable-length IPv4 header, need to use byte-based arithmetic */
    if (nh->pos + hdrsize > data_end)
        return -1;

    nh->pos += hdrsize;
    *iphdr = iph;

    return iph->protocol;
}
```

该函数的输入为三个指针，nh为一个结构体hdr_cursor的指针，该结构体中的pos对象用于追踪在xdp_md或__sk_buff结构体中当前解析的位置，data_end是指向xdp_md或__sk_buff结构体的尾部的指针，iphdr指针指向解析后的IP头部。

具体步骤如下：

1. 首先，检查nh->pos是否指向一个有效的IP报文头部，如果不是，返回-1。
这里利用了C语言的一个语法糖，结构体+1代表指针向后移动一整个结构体。如果此时的指针位置大于data_end则意味着中间不存在IP头部，返回-1。
2. 然后，计算IP报文头部的大小，因为它可能有可选字段。大小等于iph->ihl乘以四，其中iph->ihl表示报文头部的长度（以32位为单位）。
接着，检查nh->pos加上报文头部的大小是否超过了data_end，如果是，则意味着不存在有效的IP头部，返回-1。

- 最后，将`nh->pos`向后移动报文头部的大小，将`iph`赋值给`*iphdr`，返回`iph->protocol`表示IP报文的协议字段。

下一节中我们将介绍对应的eBPF程序如何调用这个解析函数。

2. 核心代码编写

本节主要以TC为例，并在注释中给出XDP的实现。

```
/* A eBPF/TC firewall to block packets from 223.5.5.5 */
int droppacket(struct __sk_buff *skb){ // XDP: int droppacket(struct
xdp_md *ctx){

    void *data_end = (void *) (long) skb->data_end;
    void *data = (void *) (long) skb->data;
    /* For XDP, just replace skb with ctx */
    struct hdr_cursor nh;
    struct ethhdr *eth;
    int eth_type;
    int ip_type;
    struct iphdr *iphdr;

    /* These keep track of the next header type and iterator pointer
    */
    nh.pos = data;
    /* Parse Ethernet */
    eth_type = parse_ethhdr(&nh, data_end, &eth);
    if (eth_type != bpf_htons(ETH_P_IP))
        return TC_ACT_OK; // XDP: return XDP_PASS;

    /* Parse IP headers */
    ip_type = parse_iphdr(&nh, data_end, &iphdr);
    if (ip_type == -1)
        return TC_ACT_OK;

    if(iphdr->saddr == 0x050505DF){ // 223.5.5.5's hex, Big-endian
mode.
        // drop the packet
        return TC_ACT_SHOT; // XDP: return XDP_DROP;
    }
    return TC_ACT_OK;
}
```

在该代码中，我们通过`nh`结构体中的`pos`对象跟踪网络数据处理位置，并依次调用以太网帧、IP报文的解析器，并根据结果返回响应。如果某个包从223.5.5.5发来，就丢弃它，反之让其通过。完整代码见本文末尾。

3. 编译与加载

在编写代码完成后，我们需要将其编译成为可以加载至系统内核的eBPF对象文件，并在终端中执行相关命令进行加载。

1. 编译

在终端中执行如下命令进行编译：

```
clang -O2 -target bpf -c tc.c -o tc.o
```

2. TC程序的加载与卸载

加载：

```
# 网卡接口需要自行通过ip a命令查看  
sudo tc qdisc add dev ens33 clsact  
sudo tc filter add dev ens33 egress bpf direct-action obj tc.o  
sec tc
```

卸载：

```
sudo tc qdisc del dev ens33 clsact
```

3. XDP程序的加载与卸载

加载：

```
# 网卡接口需要自行通过ip a命令查看  
sudo ip link set dev ens33 xdp obj xdp.o sec xdp
```

卸载：

```
sudo ip link set dev ens33 xdp off
```

五、实验结果

两个实验的结果是一致的，在此只展示eBPF/TC的结果：

```
zhaidoudou123@mbp-Parallels-ARM-Virtual-Machine:~/eBPF/packetprocessing/tc$ sudo tc filter add dev enp0s5 ingress bpf direct-action obj tc.o sec tc
zhaidoudou123@mbp-Parallels-ARM-Virtual-Machine:~/eBPF/packetprocessing/tc$ ping 225.5.5.5
PING 225.5.5.5 (225.5.5.5) 56(84) bytes of data.
^C
--- 225.5.5.5 ping statistics ---
3 packets transmitted, 0 received, 100% packet loss, time 2030ms
```

六、参考资料及建议

- XDP-tutorial
<https://github.com/xdp-project/xdp-tutorial>

七、新尝试

之前的内核模块课上的额外实验中，我们对ICMP协议以及本机发出的目标为223.5.5.5数据包进行了屏蔽，在eBPF下我们同样也可以做到这件事，只需要修改`iph->saddr`为`iph->daddr`即可。需要注意的是，对于流出方向的数据包，只有eBPF/TC可以进行相关的操作，XDP做不到。

而本周的新尝试是，实现一个针对特定端口的防火墙。

例如，ssh服务的端口号是22，且使用TCP协议。我们的目标是屏蔽一切目标端口为22的TCP报文。

提示

1. 协议的判断可以通过判断`iph->protocol`是否为`IPPROTO_TCP`实现，记得首先`#include <linux/in.h>`。
2. 先利用`parse_tcphdr()`函数解析TCP头部后再进行端口判断，该函数定义在`parsing_helpers.h`头文件中，返回值为`int`类型，值为TCP头部中储存的TCP头部长度。
3. 端口由`tcphdr`结构体的`dest`或`source`对象决定，本次只需要针对流入流量处理即可，所以选择`dest`对象，即指向本机的目标端口。具体用法为`tcphdr->dest`。另外，端口号为大段顺序的16位整数，22对应的大段十六进制端口号为`0x1600`。

八、代码附录

1. tc.c

```
/* SPDX-License-Identifier: GPL-2.0 */
#include <linux/bpf.h>
#include <linux/pkt_cls.h>
#include "bpf/bpf_helpers.h"
#include "bpf/bpf_endian.h"
#include "parsing_helpers.h"

SEC("tc")
int droppacket(struct __sk_buff *skb){
    void *data_end = (void *)(long)skb->data_end;
    void *data = (void *)(long)skb->data;
    struct hdr_cursor nh;
    struct ethhdr *eth;
    int eth_type;
    int ip_type;
    struct iphdr *iphdr;

    /* These keep track of the next header type and iterator pointer */
    nh.pos = data;
    /* Parse Ethernet */
    eth_type = parse_ethhdr(&nh, data_end, &eth);
    if (eth_type != bpf_htons(ETH_P_IP)) return TC_ACT_OK;

    /* Parse IP/IPv6 headers */
    ip_type = parse_iphdr(&nh, data_end, &iphdr);
    if (ip_type == -1) return TC_ACT_OK;
```

```

    if (iphdr->saddr == 0x050505DF){ // 223.5.5.5's hex, Big-endian
mode.
        // drop the packet
        return TC_ACT_SHOT;
    }
    return TC_ACT_OK;
}
char _license[] SEC("license") = "GPL";

```

2. xdp.c

```

/* SPDX-License-Identifier: GPL-2.0 */
#include <linux/bpf.h>
#include "bpf/bpf_helpers.h"
#include "bpf/bpf_endian.h"
#include "parsing_helpers.h"

SEC("xdp")
int droppacket(struct xdp_md *ctx){
    void *data_end = (void *) (long) ctx->data_end;
    void *data = (void *) (long) ctx->data;
    struct hdr_cursor nh;
    struct ethhdr *eth;
    int eth_type;
    int ip_type;
    struct iphdr *iphdr;

    /* These keep track of the next header type and iterator pointer
    */
    nh.pos = data;
    /* Parse Ethernet */
    eth_type = parse_ethhdr(&nh, data_end, &eth);
    if (eth_type != bpf_htons(ETH_P_IP)) return XDP_PASS;

    /* Parse IP/IPv6 headers */
    ip_type = parse_iphdr(&nh, data_end, &iphdr);
    if (ip_type == -1) return XDP_PASS;

    if (iphdr->saddr == 0x050505DF){ // 223.5.5.5's hex, Big-endian
mode.
        // drop the packet
        return XDP_DROP;
    }
    return XDP_PASS;
}
char _license[] SEC("license") = "GPL";

```

3. parsing_helpers.h

需要放在与tc.c、xdp.c相同的目录下 该文件包括了以太网、IPv4/v6、TCP/UDP、ICMP、VLAN tag 以及VXLAN等协议的解析器

```
/* SPDX-License-Identifier: (GPL-2.0-or-later OR BSD-2-clause) */
/*
 * This file contains parsing functions that are used in the packetXX
 * XDP
 * programs. The functions are marked as __always_inline, and fully
 * defined in
 * this header file to be included in the BPF program.
 *
 * Each helper parses a packet header, including doing bounds checking,
 * and
 * returns the type of its contents if successful, and -1 otherwise.
 *
 * For Ethernet and IP headers, the content type is the type of the
 * payload
 * (h_proto for Ethernet, nexthdr for IPv6), for ICMP it is the ICMP
 * type field.
 * All return values are in host byte order.
 *
 * The versions of the functions included here are slightly expanded
 * versions of
 * the functions in the packet01 lesson. For instance, the Ethernet
 * header
 * parsing has support for parsing VLAN tags.
 */

#ifndef __PARSING_HELPERS_H
#define __PARSING_HELPERS_H

#include <stddef.h>
#include <linux/if_ether.h>
#include <linux/if_packet.h>
#include <linux/ip.h>
#include <linux/ipv6.h>
#include <linux/icmp.h>
#include <linux/icmpv6.h>
#include <linux/udp.h>
#include <linux/tcp.h>

/* Header cursor to keep track of current parsing position */
struct hdr_cursor {
    void *pos;
};

/*
 * struct vlan_hdr - vlan header
 * @h_vlan_TCI: priority and VLAN ID
 * @h_vlan_encapsulated_proto: packet type ID or len
 */
```

```

struct vlan_hdr {
    __be16  h_vlan_TCI;
    __be16  h_vlan_encapsulated_proto;
};

/* VXLAN protocol (RFC 7348) header:
 * +-----+
 * |R|R|R|R|I|R|R|R|          Reserved          |
 * +-----+
 * |          VXLAN Network Identifier (VNI) |   Reserved   |
 * +-----+
 *
 * I = VXLAN Network Identifier (VNI) present.
 */
struct vxlanhdr {
    __be32 vx_flags;
    __be32 vx_vni;
};

/*
 * Struct icmphdr_common represents the common part of the icmphdr and
 * icmp6hdr
 * structures.
 */
struct icmphdr_common {
    __u8      type;
    __u8      code;
    __sum16 cksum;
};

/* Allow users of header file to redefine VLAN max depth */
#ifndef VLAN_MAX_DEPTH
#define VLAN_MAX_DEPTH 2
#endif

#define VLAN_VID_MASK      0xffff /* VLAN Identifier */
/* Struct for collecting VLANs after parsing via parse_ethhdr_vlan */
struct collect_vlans {
    __u16 id[VLAN_MAX_DEPTH];
};

static __always_inline int proto_is_vlan(__u16 h_proto)
{
    return !(h_proto == bpf_htons(ETH_P_8021Q) ||
            h_proto == bpf_htons(ETH_P_8021AD));
}

/* Notice, parse_ethhdr() will skip VLAN tags, by advancing nh->pos
and returns
 * next header EtherType, BUT the ethhdr pointer supplied still points
to the
 * Ethernet header. Thus, caller can look at eth->h_proto to see if
this was a
 * VLAN tagged packet.

```



```

*/
static __always_inline int parse_ethhdr_vlan(struct hdr_cursor *nh,
                                              void *data_end,
                                              struct ethhdr **ethhdr,
                                              struct collect_vlans *vlans)
{
    struct ethhdr *eth = nh->pos;
    int hdrsize = sizeof(*eth);
    struct vlan_hdr *vlh;
    __u16 h_proto;
    int i;

    /* Byte-count bounds check; check if current pointer + size of
header
* is after data_end.
*/
    if (nh->pos + hdrsize > data_end)
        return -1;

    nh->pos += hdrsize;
    *ethhdr = eth;
    vlh = nh->pos;
    h_proto = eth->h_proto;

    /* Use loop unrolling to avoid the verifier restriction on loops;
* support up to VLAN_MAX_DEPTH layers of VLAN encapsulation.
*/
    #pragma unroll
    for (i = 0; i < VLAN_MAX_DEPTH; i++) {
        if (!proto_is_vlan(h_proto))
            break;

        if (vlh + 1 > data_end)
            break;

        h_proto = vlh->h_vlan_encapsulated_proto;
        if (vlans) /* collect VLAN ids */
            vlans->id[i] =
                (bpf_ntohs(vlh->h_vlan_TCI) & VLAN_VID_MASK);

        vlh++;
    }

    nh->pos = vlh;
    return h_proto; /* network-byte-order */
}

static __always_inline int parse_ethhdr(struct hdr_cursor *nh,
                                          void *data_end,
                                          struct ethhdr **ethhdr)
{
    /* Expect compiler removes the code that collects VLAN ids */
    return parse_ethhdr_vlan(nh, data_end, ethhdr, NULL);
}

```

```

static __always_inline int parse_ip6hdr(struct hdr_cursor *nh,
                                       void *data_end,
                                       struct ipv6hdr **ip6hdr)
{
    struct ipv6hdr *ip6h = nh->pos;

    /* Pointer-arithmetic bounds check; pointer +1 points to after end
of
    * thing being pointed to. We will be using this style in the
remainder
    * of the tutorial.
    */
    if (ip6h + 1 > data_end)
        return -1;

    nh->pos = ip6h + 1;
    *ip6hdr = ip6h;

    return ip6h->nexthdr;
}

static __always_inline int parse_iphdr(struct hdr_cursor *nh,
                                       void *data_end,
                                       struct iphdr **iphdr)
{
    struct iphdr *iph = nh->pos;
    int hdrsize;

    if (iph + 1 > data_end)
        return -1;

    hdrsize = iph->ihl * 4;
    /* Sanity check packet field is valid */
    if(hdrsize < sizeof(*iph))
        return -1;

    /* Variable-length IPv4 header, need to use byte-based arithmetic
*/
    if (nh->pos + hdrsize > data_end)
        return -1;

    nh->pos += hdrsize;
    *iphdr = iph;

    return iph->protocol;
}

static __always_inline int parse_icmp6hdr(struct hdr_cursor *nh,
                                          void *data_end,
                                          struct icmp6hdr **icmp6hdr)
{
    struct icmp6hdr *icmp6h = nh->pos;

```

```

    if (icmp6h + 1 > data_end)
        return -1;

    nh->pos    = icmp6h + 1;
    *icmp6hdr = icmp6h;

    return icmp6h->icmp6_type;
}

static __always_inline int parse_icmphdr(struct hdr_cursor *nh,
                                         void *data_end,
                                         struct icmphdr **icmphdr)
{
    struct icmphdr *icmph = nh->pos;

    if (icmph + 1 > data_end)
        return -1;

    nh->pos    = icmph + 1;
    *icmphdr = icmph;

    return icmph->type;
}

static __always_inline int parse_icmphdr_common(struct hdr_cursor *nh,
                                                void *data_end,
                                                struct icmphdr_common **icmphdr)
{
    struct icmphdr_common *h = nh->pos;

    if (h + 1 > data_end)
        return -1;

    nh->pos    = h + 1;
    *icmphdr = h;

    return h->type;
}

/*
 * parse_vxlanhdr: parse the vxlan header and return the VNI of the udp
 * payload
 */
static __always_inline int parse_vxlanhdr(struct hdr_cursor *nh,
                                           void *data_end,
                                           struct vxlanhdr **vxlanhdr)
{
    // int len;
    struct vxlanhdr *h = nh->pos;

    if (h + 1 > data_end)
        return -1;

    nh->pos    = h + 1;

```

[illegible]

```
    if (nh->pos + len > data_end)
        return -1;

    nh->pos += len;
    *tcphdr = h;

    return len;
}

#endif /* __PARSING_HELPERS_H */
```